

# Updating Materialized Views and Caches Using Kafka

*-or-*

## Why You Should Publish Data Changes to Kafka

Zach Cox

Prairie.Code() Oct 2016

<https://github.com/zcox/twitter-microservices-example>

# About Me

- Building things with Apache Kafka since 2014
  - Currently at Uptake in Chicago: predictive analytics for industrial IoT
  - Previously at Banno in Des Moines: ad targeting for bank web sites
- Co-founded Pongr
  - Startup in Des Moines: photo marketing platform powered by messaging systems
- Software game since 1998
- Links
  - <http://theza.ch>
  - <https://github.com/zcox>
  - <https://twitter.com/zcox>
  - <https://www.linkedin.com/in/zachcox>

# Remember These Things

1. Learn about Apache Kafka <http://kafka.apache.org>
2. Send events and data changes to Kafka
3. Denormalization is OK
4. **Up-to-date materialized views and caches**

# Build a New Service

- Provides read access to data from many sources
  - Query multiple tables or databases
  - Complex joins, aggregations
- Response latency
  - 95% 5 msec
  - Max 10 msec
- Data update latency
  - 95% 1 sec
  - Max 10 sec



**Zach Cox**

@ZCOX

Engineering @Uptake

📍 Chicago, IL

🔗 theza.ch

📅 Joined January 2008

📷 97 Photos and videos



TWEETS  
3,815

FOLLOWING  
1,093

FOLLOWERS  
894

LIKES  
183

LISTS  
1

MOMENTS  
0

Tweets

Tweets & replies

Media



**Zach Cox** @zcox · 13m

So unprofessional when conference speakers are still writing slides the day of their talk. Except when it's me, then it's totally fine...





Home



Moments



Notifications



Messages



Search Twitter

**Zach Cox**

@ZCOX

TWEETS

3,815

FOLLOWING

1,093

FOLLOWERS

894

## Trends · Change

## #Titanfall

Titanfall 2 Coming Friday to Xbox One, PlayStation 4 and Origin for PC

Promoted by Titanfall

**Anthony Davis**

87.3K Tweets

**#LakeShow**

95K Tweets



What's happening?

**Skills Matter** @skillsmatter · 11sThe @LifeatIG are hosting a #droidconUK trading comp! Win a brand new Google Pixel 128 GB phone! More info: [buff.ly/2e3Ns6M](http://buff.ly/2e3Ns6M) @IGlabs**IG trading competition**

IG trading competition

[ig.com](http://ig.com)

In reply to Randy Bias

**adrian cockcroft** @adrianco · 19s

.@randybias so coding for portability first is often premature optimization.





Edit profile

**Zach Cox**

@ZCOX

Engineering @Uptake

Chicago, IL [theza.ch](https://theza.ch)

1,093 FOLLOWING

894 FOLLOWERS

TWEETS

MEDIA

LIKES



**Zach Cox** @zcox

13s

Replying to Kris Nuttycombe

Thanks for the inspiration! The content will at least be very fresh in my mind.



**Zach Cox** @zcox

11m

# Response Times: The 3 Important Limits

by **JAKOB NIELSEN** on January 1, 1993

Topics: [Applications](#) [Technology](#) [User Behavior](#) [Web Usability](#)

**Summary:** There are 3 main time limits (which are determined by human perceptual abilities) to keep in mind when optimizing web and application performance.

*Excerpt from Chapter 5 in my book [Usability Engineering](#), from 1993:*

The basic advice regarding response times has been about the same for thirty years [Miller 1968; Card et al. 1991]:

- **0.1 second** is about the limit for having the user feel that the system is **reacting instantaneously**, meaning that no special feedback is necessary except to display the result.
- **1.0 second** is about the limit for the **user's flow of thought** to stay uninterrupted, even though the user will notice the delay. Normally, no special feedback is necessary during delays of more than 0.1 but less than 1.0 second, but the user does lose the feeling of operating directly on the data.
- **10 seconds** is about the limit for **keeping the user's attention** focused on the dialogue. For longer delays, users will want to perform other tasks while waiting for the computer to finish, so they should be given feedback indicating when the computer expects to be done. Feedback during the delay is especially important if the response time is likely to be highly variable, since users will then not know what to expect.

<http://www.nngroup.com/articles/response-times-3-important-limits/>



# User Information Service

- One operation: get user information
  - REST HTTP + JSON
  - Input: userId
  - GET /users/:userId
- Output:
    - userId
    - username
    - Name
    - Description
    - Location
    - Web page url
    - Joined date
    - Profile image url
    - Background image url
    - # tweets
    - # following
    - # followers
    - # likes
    - # lists
    - # moments

# Existing RDBMS with Normalized Tables

- users
  - user\_id
  - username
  - name
  - description
- tweets
  - tweet\_id
  - text
  - user\_id (FK users)
- follows
  - follow\_id
  - follower\_id (FK users)
  - followee\_id (FK users)
- likes
  - like\_id
  - user\_id (FK users)
  - tweet\_id (FK tweets)

# Standard Solution: Query Existing Tables

- **User fields**

- `SELECT * FROM users WHERE user_id = ?`

- **# tweets**

- `SELECT COUNT(*) FROM tweets WHERE user_id = ?`

- **# following**

- `SELECT COUNT(*) FROM follows WHERE follower_id = ?`

- **# followers**

- `SELECT COUNT(*) FROM follows WHERE followee_id = ?`

- **# likes**

- `SELECT COUNT(*) FROM likes WHERE user_id = ?`

# Problems with Standard Solution

- Complex: multiple queries across multiple tables
- Potentially large aggregations at query time
  - Puts load on DB
  - Increases service response latency
  - Repeated on every query for same userId
- Shared data storage
  - Some other service writes to these tables (i.e. owns them)
  - When it changes schema, our service could break

# Standard Solution: Add a Cache

- e.g. Redis
- Benefits
  - Faster key lookups than RDBMS queries
  - Store expensive computed values in cache and reuse them (i.e. **materialized view**)
- Usage
  - Read from cache first, if found then return cached data
  - Otherwise, read from DB, write to cache, return cached data

```
def getUser(id: String): User =  
  readUserFromCache(id) match {  
    case Some(user) => user  
    case None =>  
      val user = readUserFromDatabase(id)  
      writeUserToCache(user)  
      user  
  }
```

```
def getUser(id: String): User =  
  readUserFromCache(id) match {  
    case Some(user) => user  
    case None => //cache miss!  
      val user = readUserFromDatabase(id)  
      writeUserToCache(user)  
      user  
  }
```

```
def getUser(id: String): User =  
  readUserFromCache(id) match {  
    case Some(user) => user //stale?  
    case None => //cache miss!  
      val user = readUserFromDatabase(id)  
      writeUserToCache(user)  
      user  
  }
```



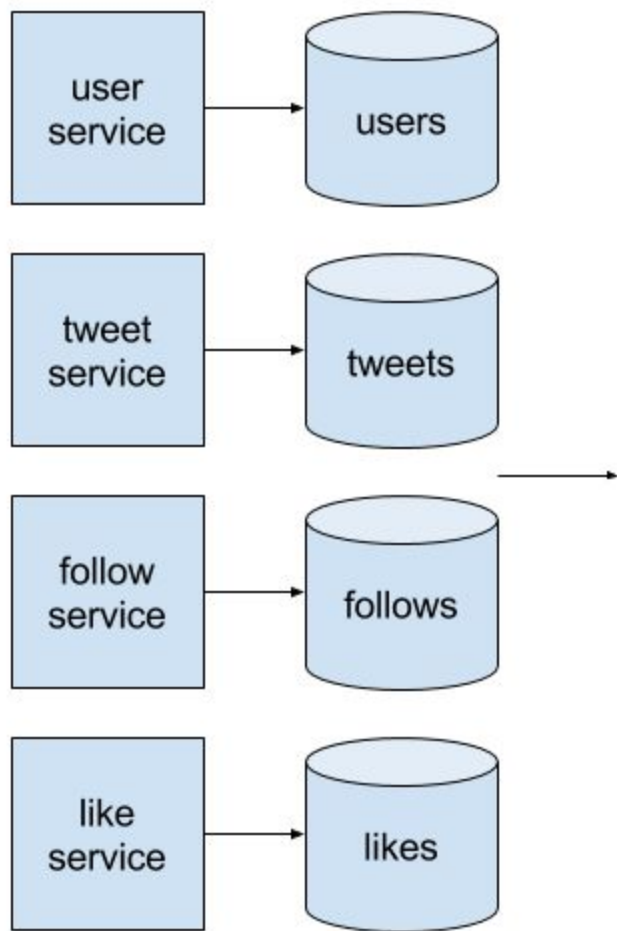
```
def getUser(id: String): User =  
  readUserFromCache(id) match { //network latency  
    case Some(user) => user //stale?  
    case None => //cache miss!  
      val user = readUserFromDatabase(id)  
      writeUserToCache(user)  
      user  
  }
```

# Problems with Standard Approach to Caches


- Operational complexity: someone has to manage Redis
- Code complexity: now querying two data stores and writing to one
- Cache misses: still putting some load on DB
- Stale data: cache is not updated when data changes
- Network latency: cache is remote

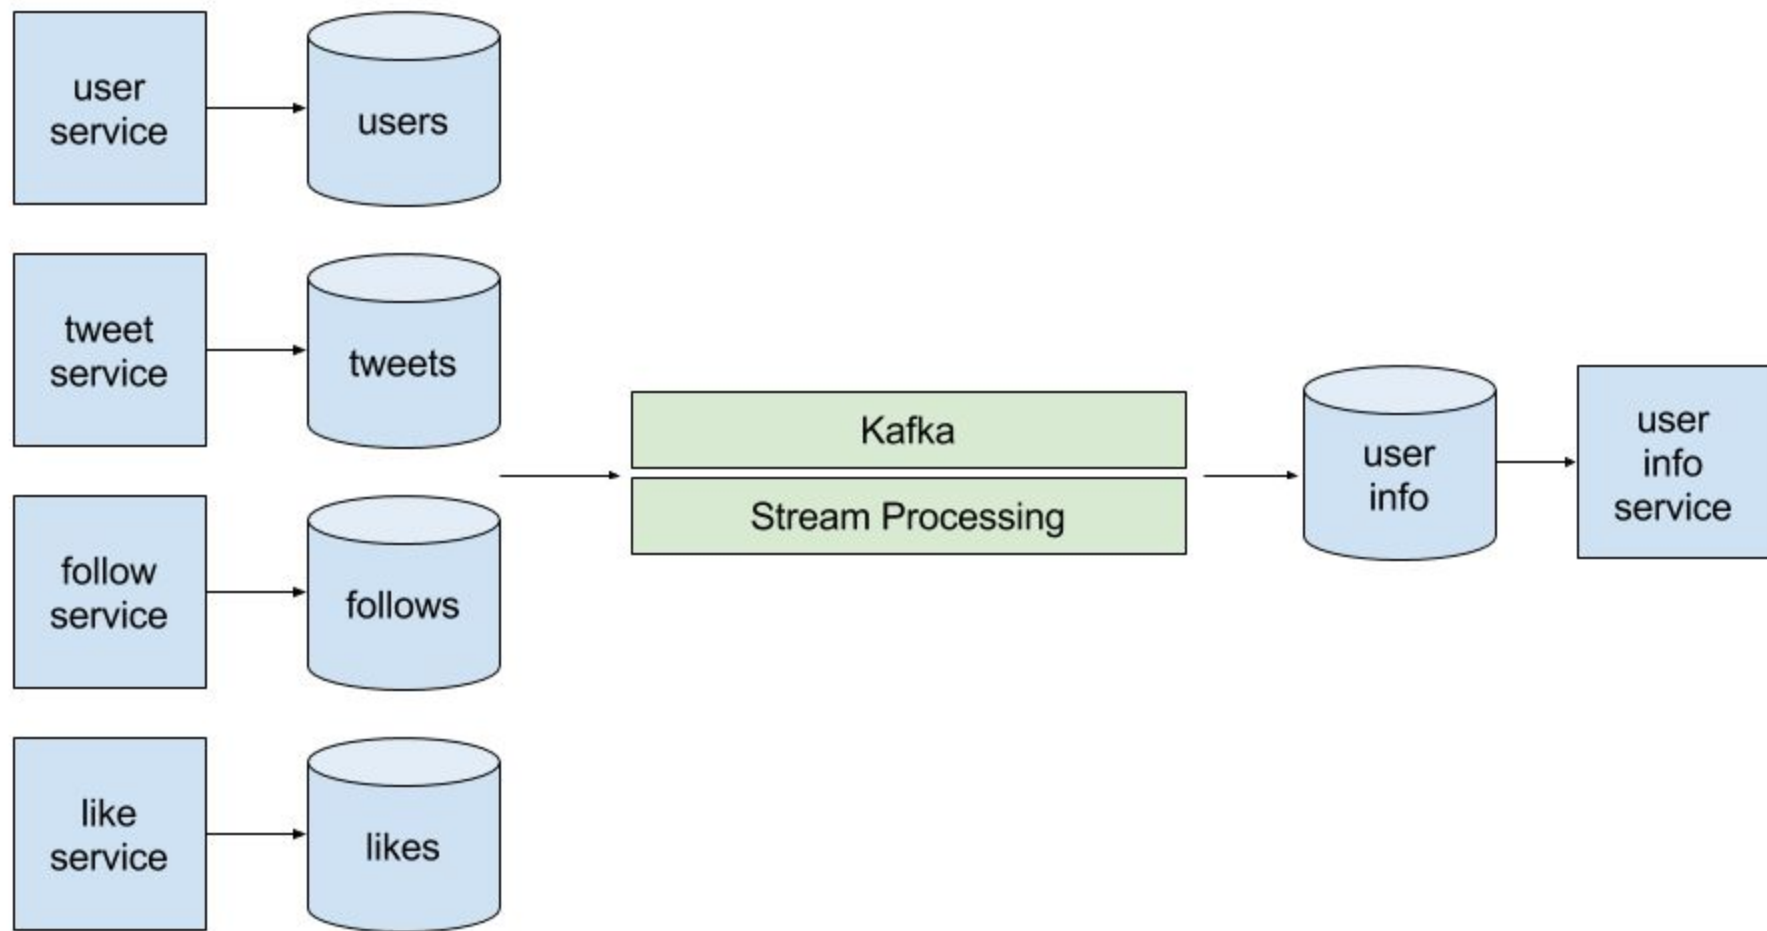
# Can We Solve These Problems?

- Yes: **If cache is always updated**
- Complexity: only read from cache
- Cache misses: cache always has *all* data
- Stale data: cache always has *updated* data
- Network latency: if cache is local to service (bonus)




?





# Kafka: Topics, Producers, Consumers



- Horizontally scalable, durable, highly available, high throughput, low latency

# Kafka: Messages

- Message is a (key, value) pair
- Key and value are byte arrays (BYO serialization)
- Key is typically an ID (e.g. userId)
- Value is some payload (e.g. page view event, user data updated)

# Kafka: Producer API

```
val props = ... //kafka host:port, other configs
```

```
val producer = new KafkaProducer[K, V](props)
```

```
producer.send(topic, key, value)
```



# Kafka: Consumer API

```
val props = ... //kafka host:port, other configs

val consumer = new KafkaConsumer[K, V](props)

consumer.subscribe(topics)

while (true) {


    val messages = consumer.poll(timeout)

    //process list of messages


}
```

# Kafka: Types of Topics


- Record topic
  - Finite topic retention period (e.g. 7 days)
  - Good for user activity, logs, metrics




- Changelog topic
  - Log-compacted topic: retains newest message for each key
  - Good for entities/table data




# Kafka: Tables and Changelogs are Dual



# Database Replication



**Primary Backup**



**State Machine Replication**

# DB to Kafka

- Change data capture
  - Kafka Connect <http://kafka.apache.org/documentation#connect>
  - Bottled Water <https://github.com/confluentinc/bottledwater-pg>
- Dual writes
  - Application writes to both DB and Kafka
  - Prefer CDC

# Kafka Streams

- Higher-level API than producers and consumers
- Just a library (no Hadoop/Spark/Flink cluster to maintain)

```
val tweetCountsByUserId = builder.stream(tweetsTopic)

    .selectKey((tweetId, tweet) => tweet.userId)

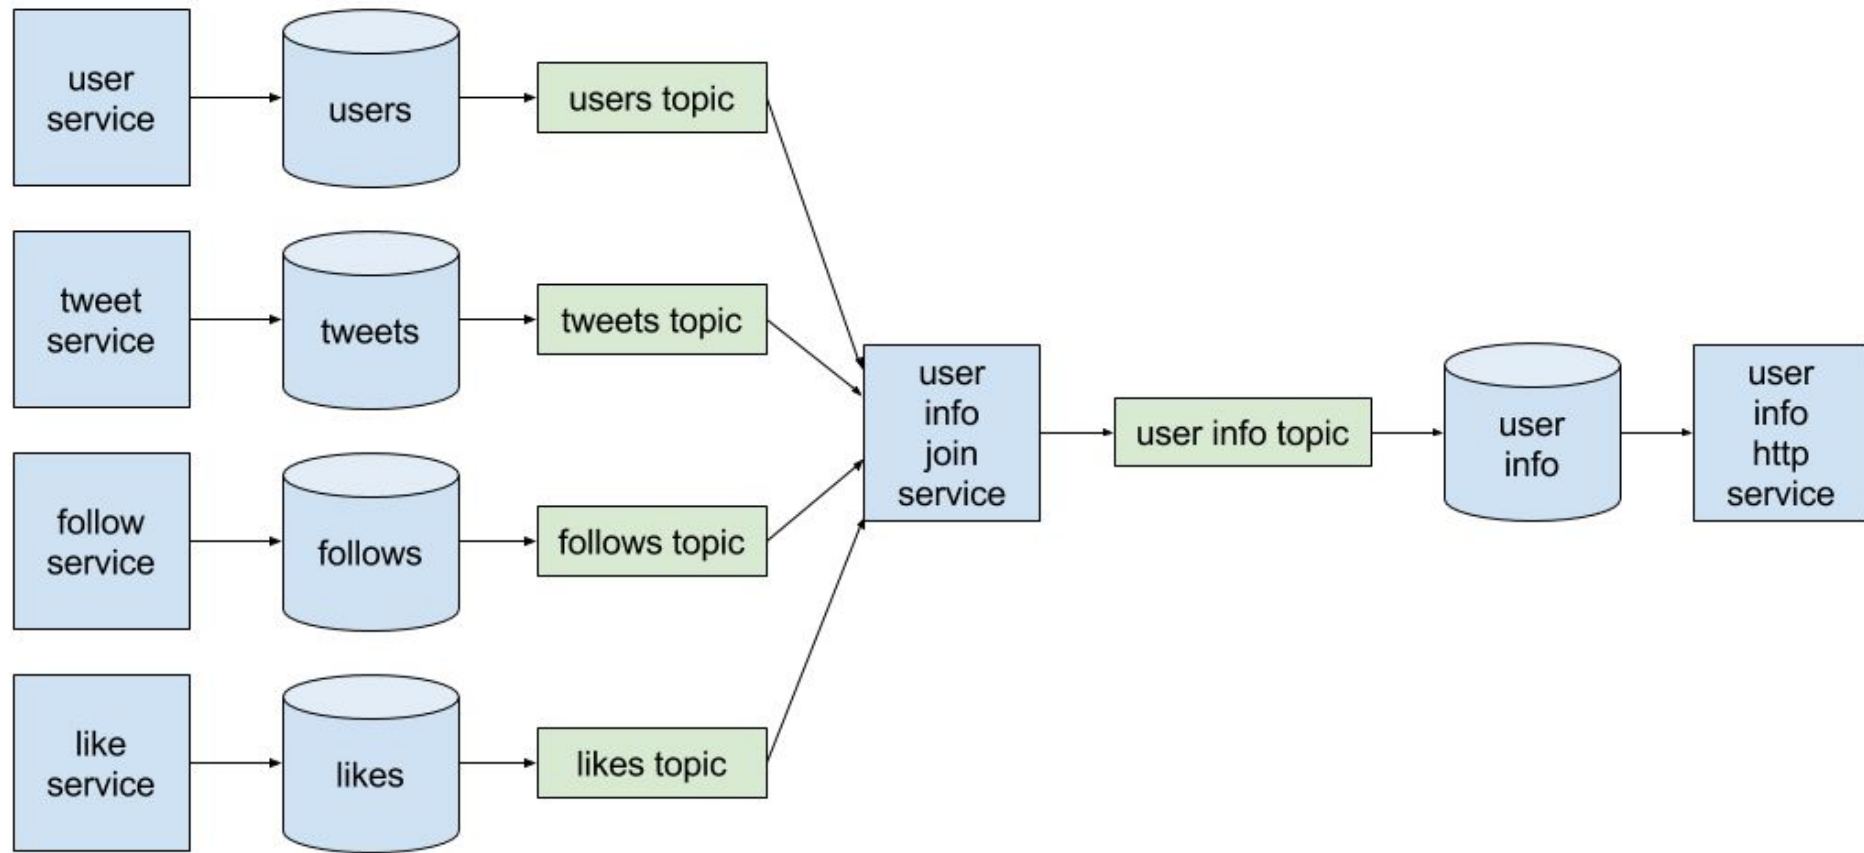
    .countByKey("tweetCountsByUserId")

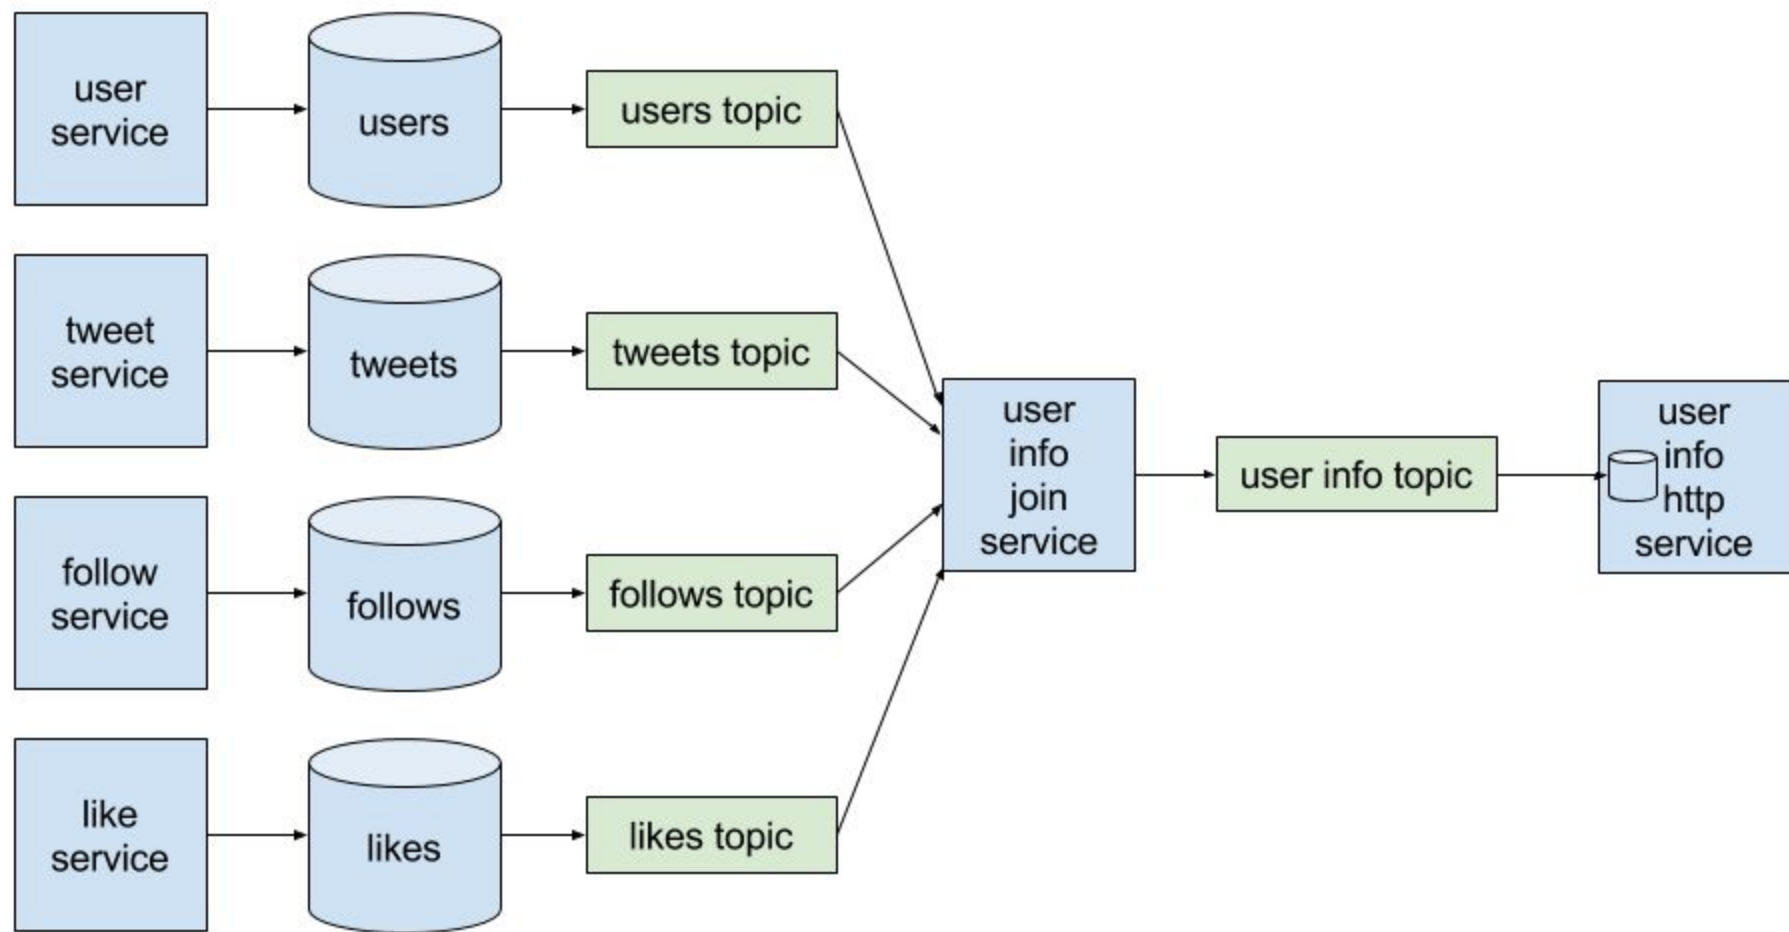
val userInformation = builder.table(usersTopic)

    .leftJoin(tweetCountsByUserId,

        (user, count) => new UserInformation(user, count))

userInformation.to(userInformationTopic)
```



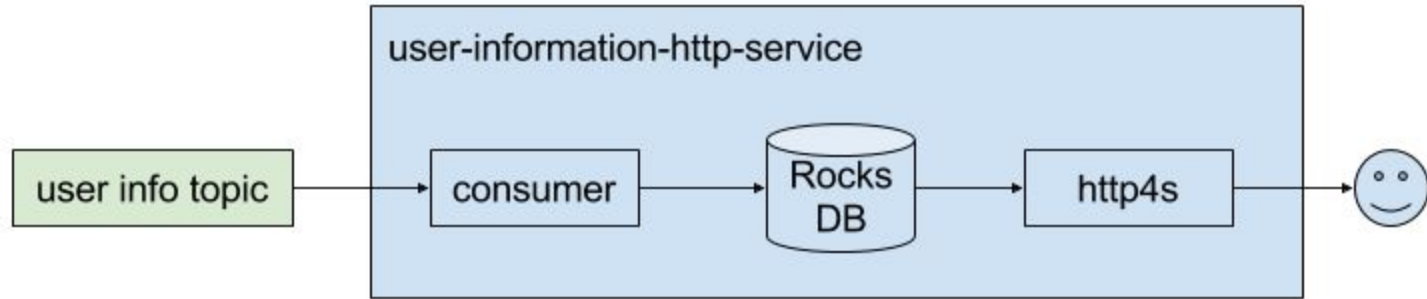




# RocksDB

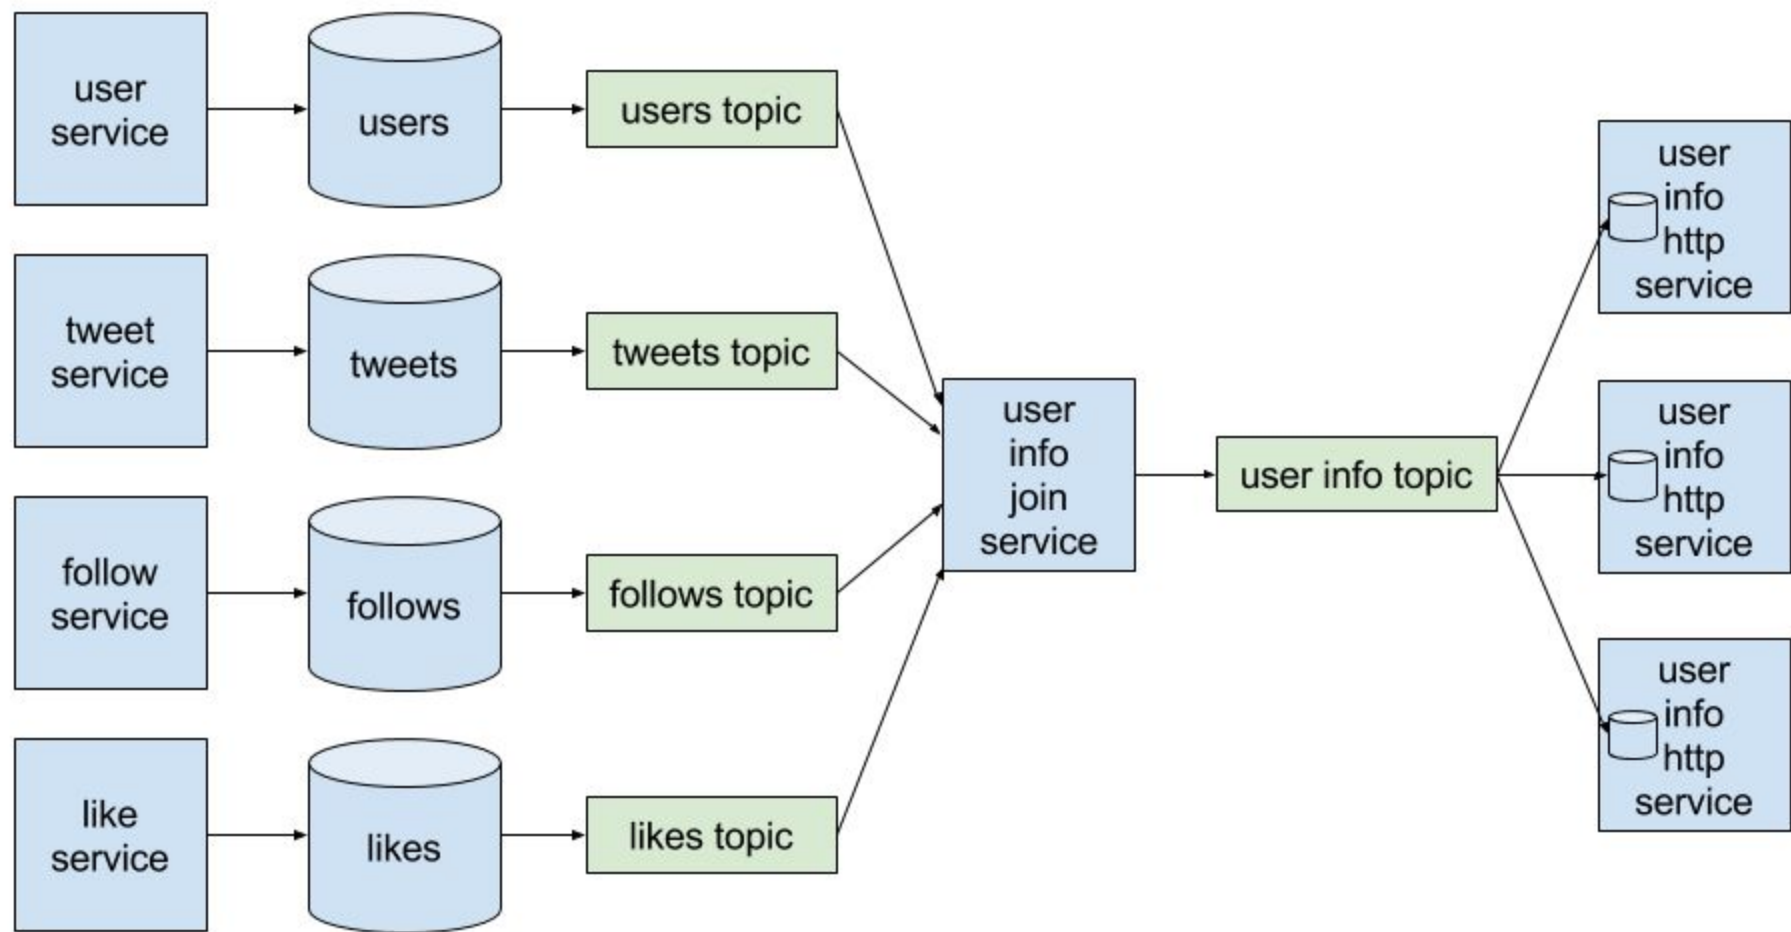
- Key-value store
- In-process
  - Local (not remote)
  - Library (not a daemon/server)
- Mostly in-memory, spills to local disk
  - Usually an under-utilized resource on app servers
  - 100s of GBs? TBs?
  - AWS EBS 100GB SSD \$10/mo
- <http://rocksdb.org>

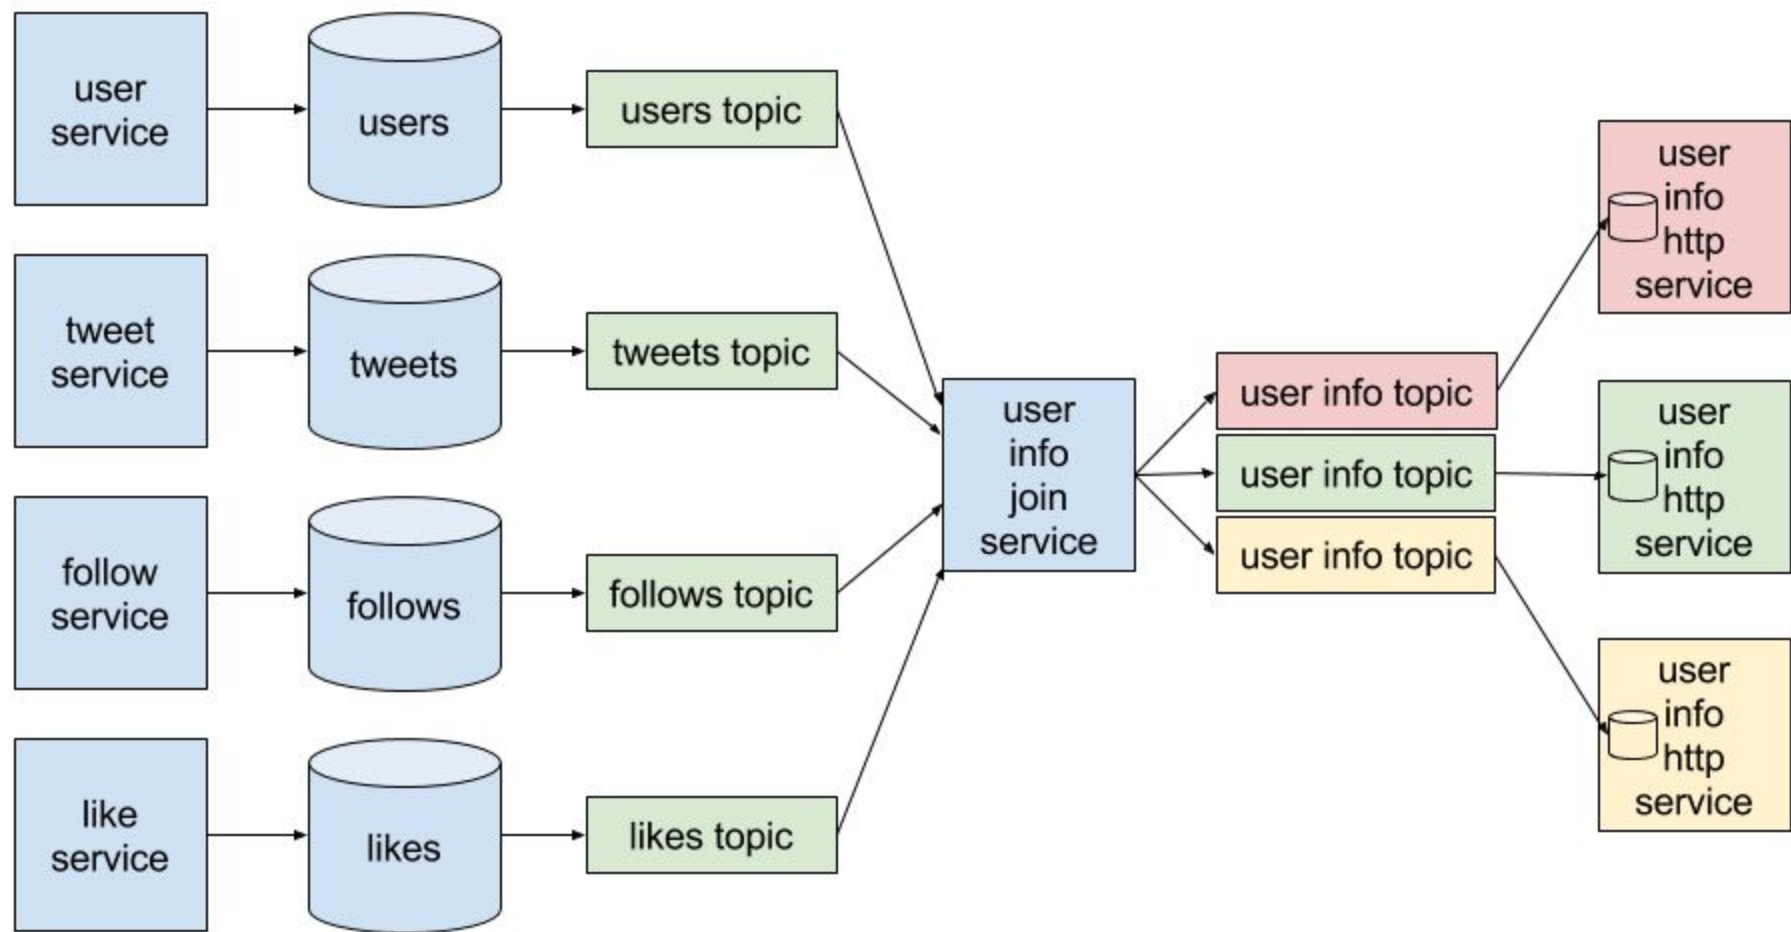
# HTTP Service Internals

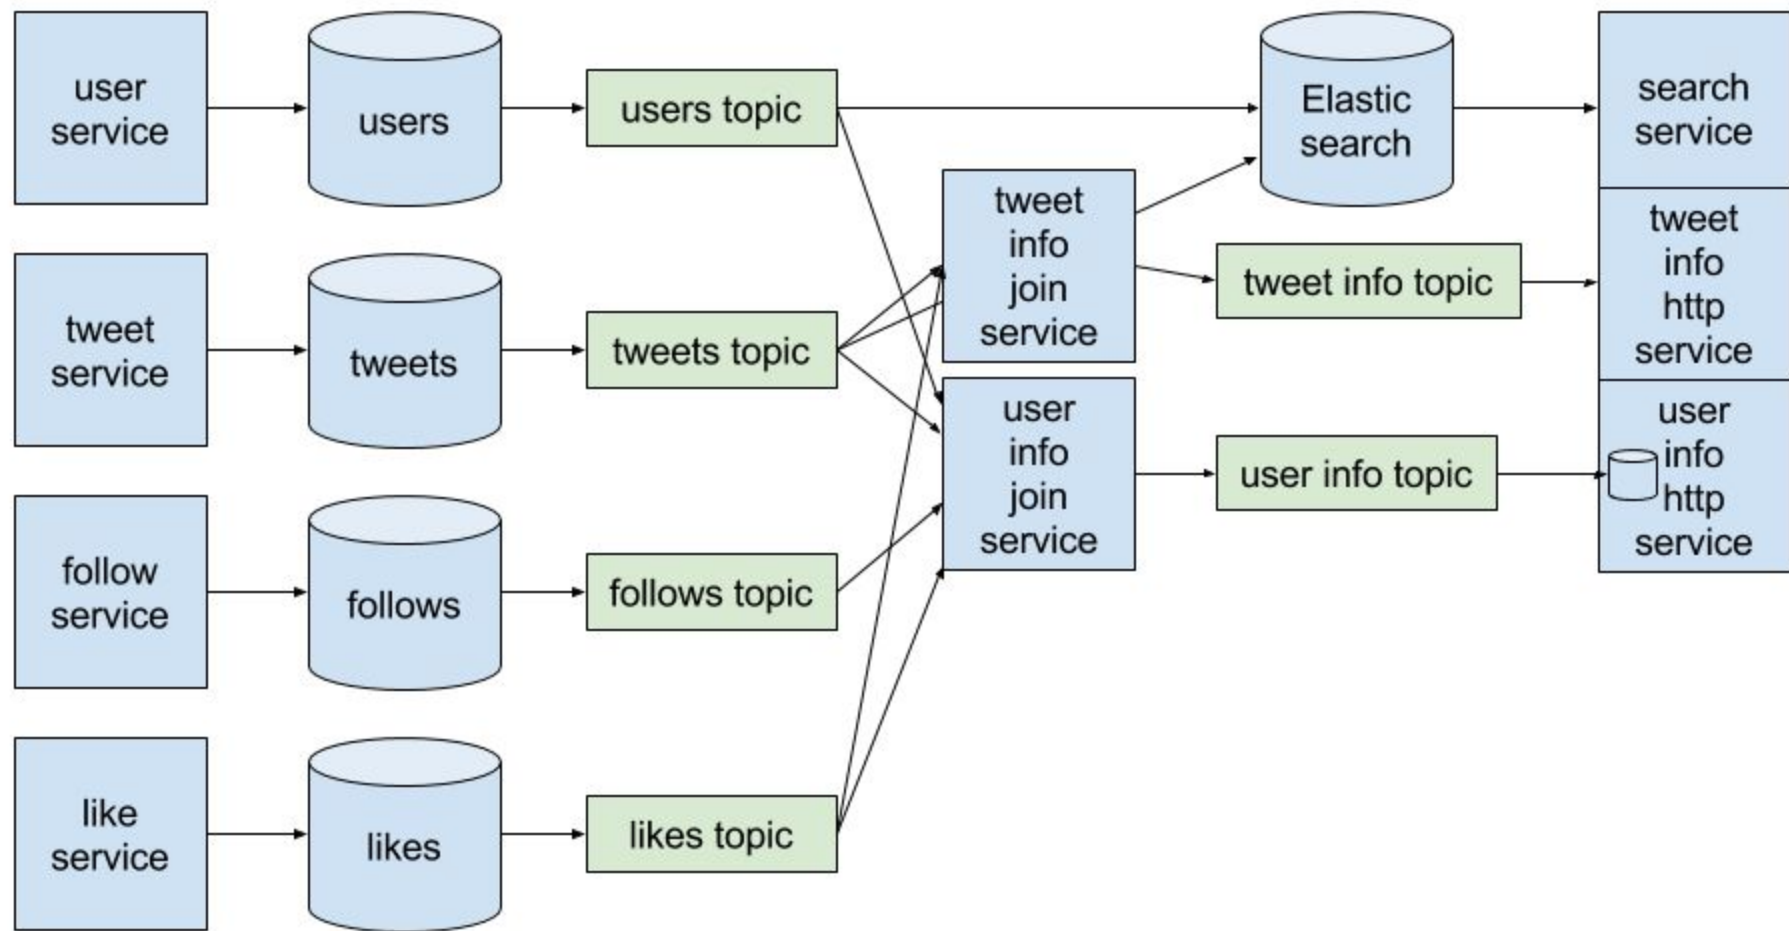


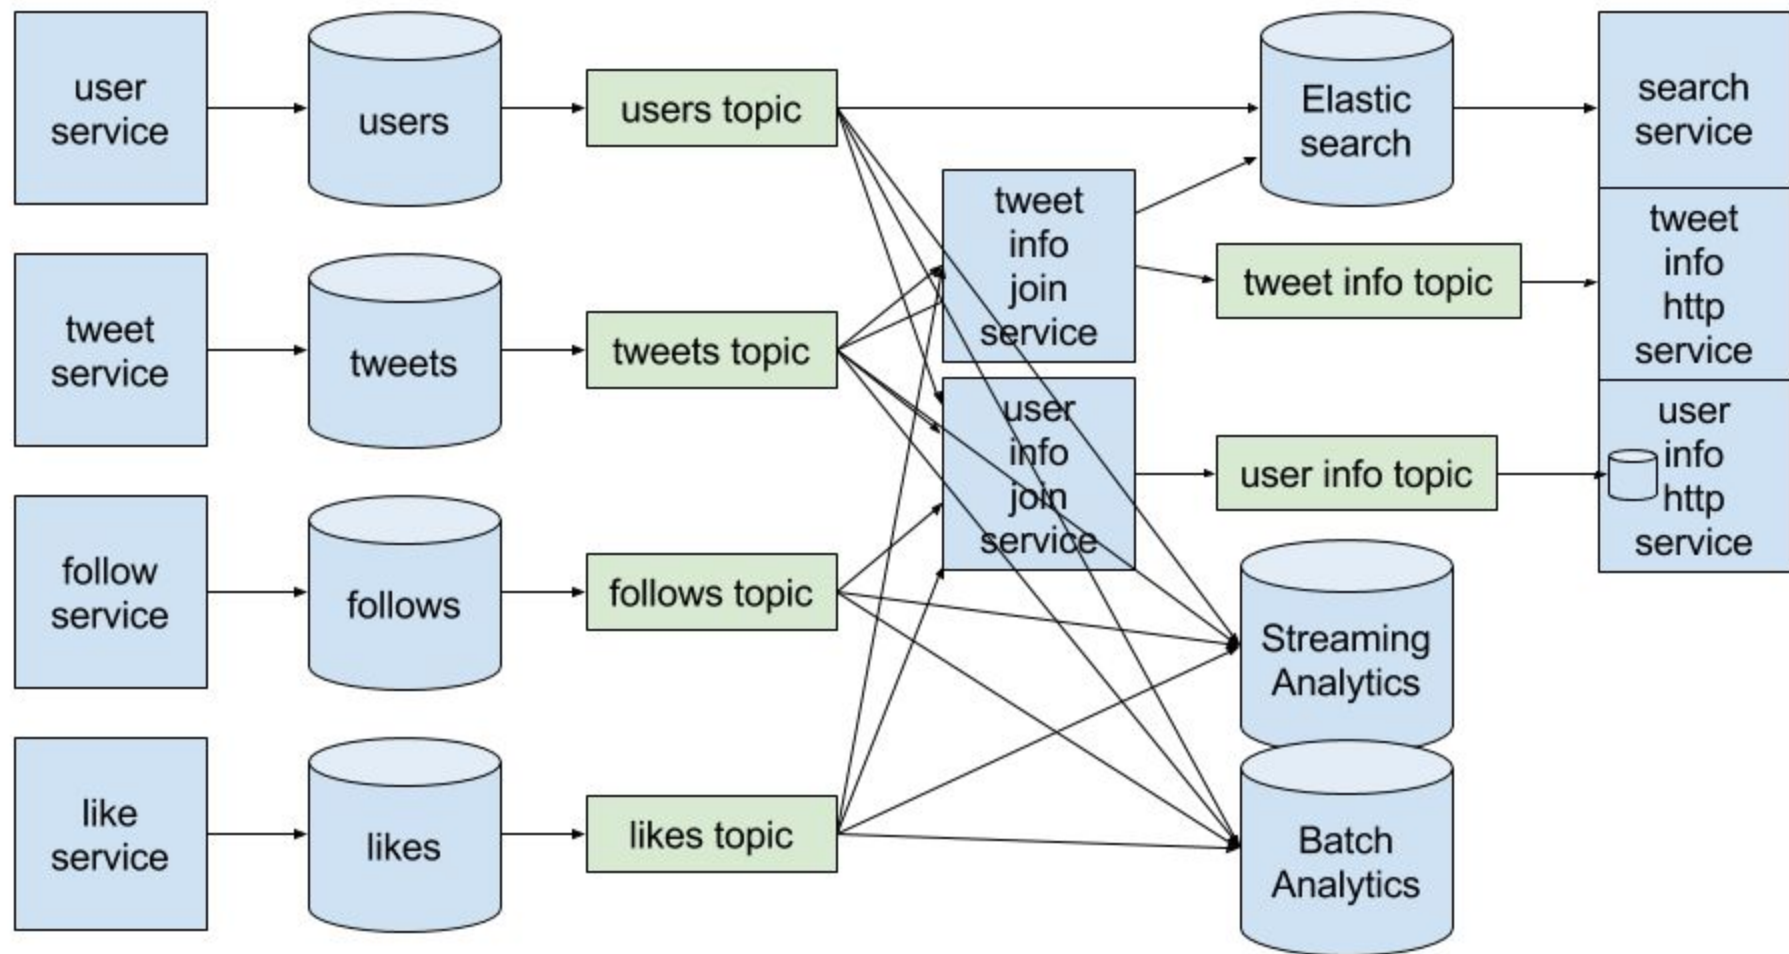
# Live Demo!











# Kafka Streams Interactive Queries




Jay Kreps

@jaykrep



Following

1. What does it mean to have "interactive query" support for Kafka's Streams API? The long story is in the blog:



## Unifying Stream Processing and Interactive Queries in Apa...

We are excited to announce Interactive Queries, a new feature for stream processing with Apache Kafka™.

[confluent.io](http://confluent.io)

RETWEETS

32

LIKES

64



6:21 PM - 26 Oct 2016

<http://www.confluent.io/blog/unifying-stream-processing-and-interactive-queries-in-apache-kafka/>



# Confluent Schema Registry

- Serialize messages in Kafka topics using Avro
- Avro schemas registered with central server
- Anyone can safely consume data in topics
- <http://docs.confluent.io/3.0.1/schema-registry/docs/index.html>